

# Federated Learning for Industrial Anomaly Detection

Stage 1: Baseline Development & Independent Client Setup

AI for Trustworthy Decision Making

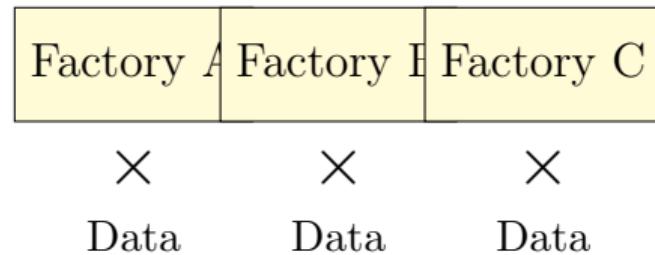
AutoVI Dataset Project

November 26, 2025

# The Challenge: Privacy in Industrial Quality Control

## Current Problem:

- ▶ Manufacturing data is **proprietary**
- ▶ Different factories = **isolated data**
- ▶ Regulatory **restrictions** on data sharing
- ▶ Cannot train shared models easily



---

## Solution: Federated Learning

- ✓ Train models **locally**, share only **model updates**
- ✓ No raw data leaves the facility

# Dataset: Automotive Visual Inspection (AutoVI)

Real industrial data from Renault Group

## 6 Product Categories:

- ▶ engine\_wiring (285 train)
- ▶ pipe\_clip (195 train)
- ▶ pipe\_staple (191 train)
- ▶ tank\_screw (318 train)
- ▶ underbody\_pipes (161 train)
- ▶ underbody\_screw (373 train)

## Dataset Statistics:

Metric	Value
Total Images	3,950
Training Images	1,523
Test Images	2,399
Categories	6
Defect Types	10

% TODO: Insert image grid (6 categories)

- ▶ **Unsupervised setup:** Train on “good” images only
- ▶ **Real world:** Lighting variation, authentic defects
- ▶ **Binary annotations:** Ground truth masks for localization

# Baseline Model: PatchCore Architecture

## Memory Bank-Based Anomaly Detection

### 1. Feature Extraction

- ▶ Pre-trained WideResNet-50-2 backbone
- ▶ Multi-scale features (Layer 2 + Layer 3)
- ▶ Output: 1536-dim patch embeddings

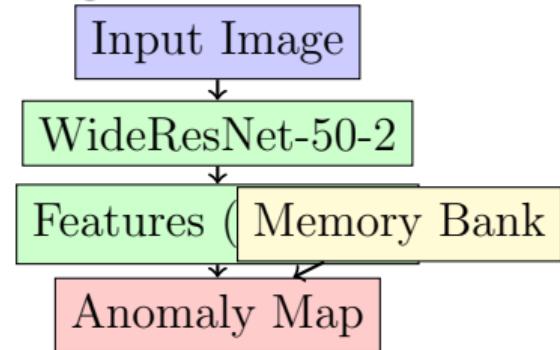
### 2. Memory Bank (Coreset)

- ▶ Greedy selection (10% of patches)
- ▶ Represents “normal” feature space
- ▶ Computationally efficient

### 3. Anomaly Scoring

- ▶ Distance to nearest bank feature
- ▶ Upsampled to pixel resolution

% TODO: Architecture diagram



# Federated Architecture: 5 Clients with IID Partitioning

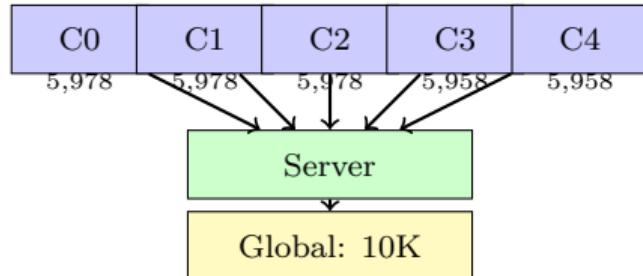
## Balanced data distribution across clients

Client	Samples	Patches	Coreset (10%)
Client 0	305	59,780	5,978
Client 1	305	59,780	5,978
Client 2	305	59,780	5,978
Client 3	304	59,584	5,958
Client 4	304	59,584	5,958
<b>Total</b>	<b>1,523</b>	<b>298,508</b>	<b>29,850</b>

**IID Partitioning:** Each client receives samples from all 6 categories

- ▶ **Balanced distribution:** ~305 samples per client
- ▶ Each client has **mixed categories** (engine\_wiring, pipe\_clip, etc.)
- ▶ Local coreset selection: **10%** of extracted patches

# Federated Memory Bank Aggregation



## Local Processing:

- ▶ Extract 784 patches/image
- ▶ Coreset: 10% selection
- ▶ Total: 29,850 patches

## Server Aggregation:

- ▶ Weighted 2x oversampling
- ▶ Diversity selection
- ▶ Final: 10,000 patches

**Compression:** 298K → 30K → 10K (**96.6% reduction**)

# Aggregation Strategies

Strategy	Description	Trade-off	Used
<b>Federated Coreset</b>	Weighted + oversampling + diversity	Fairness + diversity	✓
Simple Concatenate Diversity Preserving	Direct merge Min. 100 patches/client	Fast, no fairness Small client voice	

## Federated Coreset Algorithm:

1. Weight contributions by sample count
2. Oversample 2x from each client
3. Apply greedy coresset selection

**Client Weights (IID):**  $\sim 0.20$  each

*Only embeddings shared – raw images never leave clients*

# Experimental Setup & Training Statistics

## Configuration:

- ▶ Clients: 5 (IID)
- ▶ Backbone: WideResNet50-2
- ▶ Features: layer2 + layer3
- ▶ Dimension: 1536
- ▶ Image:  $224 \times 224$

## Results:

- ▶ Time: 173 seconds
- ▶ Rounds: 1 (one-shot)
- ▶ Local coresset: 10%
- ▶ Global bank: 10,000
- ▶ Strategy: federated\_coreset

## Data Flow:

Stage	Patches	Kept
Raw extraction	298,508	100%
Local coreset	29,850	10%
Global coreset	10,000	3.4%

# Evaluation Metrics: AUC-sPRO and AUC-ROC

## AUC-sPRO (Localization)

- ▶ Pixel-level accuracy
- ▶ Multiple FPR thresholds:
  - ▶ @0.01 (strict)
  - ▶ @0.05 (intermediate)
  - ▶ @0.1 (moderate)
  - ▶ @0.3 (permissive)
- ▶ Saturated to prevent over-crediting

## AUC-ROC (Classification)

- ▶ Image-level detection
- ▶ Binary: Good vs Anomalous
- ▶ Standard metric
- ▶ Range: [0, 1]
- ▶ Higher is better

## % TODO: Figure – FPR-sPRO curves for different categories

Expected visualization: Multiple curves comparing categories

- ▶ Performance varies by product complexity
- ▶ Starts at 1 before reaching the desired sensitivity

# Stage 1 Results: Federated Training Complete

## Training Statistics (Completed):

Client	Samples	Patches	Coreset	Weight
Client 0	305	59,780	5,978	0.200
Client 1	305	59,780	5,978	0.200
Client 2	305	59,780	5,978	0.200
Client 3	304	59,584	5,958	0.200
Client 4	304	59,584	5,958	0.200
<b>Total/Avg</b>	<b>1,523</b>	<b>298,508</b>	<b>29,850</b>	<b>1.000</b>

## Evaluation Metrics (TODO):

Category	@FPR=0.01	@FPR=0.05	@FPR=0.1	@FPR=0.3	AUC-ROC
engine_wiring	TODO	TODO	TODO	TODO	TODO
pipe_clip	TODO	TODO	TODO	TODO	TODO
pipe_staple	TODO	TODO	TODO	TODO	TODO
tank_screw	TODO	TODO	TODO	TODO	TODO
underbody_pipes	TODO	TODO	TODO	TODO	TODO
underbody_screw	TODO	TODO	TODO	TODO	TODO

**Key Achievement:** Global memory bank (10,000 patches) trained in

# Stage 1 Implementation Status

Component	Status	Notes
Data Loader	✓ Complete	All 6 categories loaded
Data Partitioning	✓ Complete	5 clients, IID distribution
PatchCore Model	✓ Complete	WideResNet50-2 backbone
Federated Training	✓ Complete	173s, global bank ready
Server Aggregation	✓ Complete	federated_coreset strategy
Baseline Evaluation	In Progress	AUC metrics pending

## Completed Outputs:

- ▶ Global memory bank: 10,000 patches (1536-dim)
- ▶ Training logs and statistics saved
- ▶ Checkpoint for reproducibility

## Next Steps:

# Stage 2 Preview: Trust-Focused Enhancements

Building on Stage 1 baselines...

## Privacy Enhancement

- ▶ Differential Privacy (DP-SGD)
- ▶ Formal privacy guarantees
- ▶ Privacy budgets:  $\varepsilon = 1, 5, 10$
- ▶ Measure privacy-utility trade-off

## Fairness Enhancement

- ▶ Address data imbalance
- ▶ Reduce performance variance
- ▶ Cross-category equity
- ▶ Client contribution weighting

## Aggregation Strategy

- ▶ Memory bank pooling (1 communication round)
- ▶ Weighted coreset selection
- ▶ Fairness-aware weighting

## Analysis

- ▶ Statistical significance testing
- ▶ Trade-off Pareto frontiers
- ▶ Recommendations by use case

# Key Takeaways & Project Roadmap

## Stage 1 Achievements:

- ▶ ✓ Complete data infrastructure (6 categories, 3,950 images)
- ▶ ✓ Federated architecture (5 clients, IID partitioning)
- ▶ ✓ PatchCore model with federated aggregation complete
- ▶ ✓ Global memory bank trained (10,000 patches, 173s)

## Stage 2 Objectives:

- ▶ Add privacy guarantees (Differential Privacy)
- ▶ Implement fairness mechanisms
- ▶ Demonstrate one-round efficient aggregation
- ▶ Provide trade-off analysis

## Final Deliverables:

- ▶ Technical Report (18-20 pages)
- ▶ Complete Code Repository

# Questions?

Dataset: [doi.org/10.5281/zenodo.10459003](https://doi.org/10.5281/zenodo.10459003)

Code: [GitHub Repository]

Contact information for team members

# Backup: Detailed Dataset Statistics

Category	Train	Test Total	Test Good	Test Anom	Defect Types	Size
engine_wiring	285	607	285	322	4	400×400
pipe_clip	195	337	195	142	2	400×400
pipe_staple	191	305	188	117	1	400×400
tank_screw	318	413	318	95	1	1000×750
underbody_pipes	161	345	161	184	3	1000×750
underbody_screw	373	392	374	18	1	1000×750
<b>Total</b>	<b>1,523</b>	<b>2,399</b>	<b>1,521</b>	<b>878</b>	<b>10</b>	-

- ▶ Small images (400×400): 671 train, 1,249 test
- ▶ Large images (1000×750): 852 train, 1,150 test
- ▶ Total defect types: 10 (mix of structural and logical)

# Backup: PatchCore Algorithm Details

## Greedy Coreset Selection:

1. Start with all patches extracted from training images
2. Randomly select first patch
3. Iteratively add patch maximizing minimum distance to selected set
4. Continue until target size (10% of total patches)

## Memory Bank Construction:

- ▶ Coreset represents normal feature distribution
- ▶ Size: 10% of all extracted patches
- ▶ Stored as feature vectors (1536-D)
- ▶ Enables efficient nearest-neighbor search (FAISS)

## Inference:

- ▶ Extract patches from test image

# Backup: Federated Aggregation Strategy (Stage 2)

## Why Memory Bank Aggregation?

- ▶ Unlike gradient FL: Only 1 communication round needed
- ▶ Not iterative: Feature banks, not parameters
- ▶ Efficient: Total 450 MB communication

## Weighted Coreset Selection:

1. Weight contributions by local dataset size (fairness)
2. Oversample from each client ( $2 \times$  target allocation)
3. Apply global greedy coresset selection
4. Ensures diverse representation + balance

## Client Imbalance Handling:

- ▶ Small clients: pipe\_clip (195 images)
- ▶ Large clients: underbody\_screw (373 images)